

```
In [1]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly re mount, call drive.mount("/content/drive", force_remount=True).

Import packages

```
In [0]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import nltk
import re
import pickle
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
from sklearn.metrics import classification_report, confusion_matrix
, accuracy_score
```

Dataset propocessing

1.Read the dataset

```
In [0]: path = "drive/My Drive/Eluvio/Eluvio_DS_Challenge.csv"
df = pd.read_csv(path)
```

```
In [4]: df.head()
```

```
Out[4]:
```

	time_created	date_created	up_votes	down_votes	title	over_18	author	cate
0	1201232046	2008-01-25	3	0	Scores killed in Pakistan clashes	False	polar	world
1	1201232075	2008-01-25	2	0	Japan resumes refuelling mission	False	polar	world
2	1201232523	2008-01-25	3	0	US presses Egypt on Gaza border	False	polar	world
3	1201233290	2008-01-25	1	0	Jump-start economy: Give health care to all	False	fadi420	world
4	1201274720	2008-01-25	4	0	Council of Europe bashes EU&UN terror blacklist	False	mhermans	world

```
In [5]: len(df)
```

```
Out[5]: 509236
```

```
In [6]: print(sum(df['category'] == "worldnews"))
print(sum(df["down_votes"] == 0))
```

```
509236
```

```
509236
```

All category is "worldnews" and all "down_votes" are 0, so dropped

```
In [0]: df = df.drop("category", axis = 1)
df = df.drop("down_votes", axis = 1)
df = df.drop("time_created", axis = 1)
df = df.drop("date_created", axis = 1)
```

```
In [8]: df.head()
```

```
Out[8]:
```

	up_votes	title	over_18	author
0	3	Scores killed in Pakistan clashes	False	polar
1	2	Japan resumes refuelling mission	False	polar
2	3	US presses Egypt on Gaza border	False	polar
3	1	Jump-start economy: Give health care to all	False	fadi420
4	4	Council of Europe bashes EU&UN terror blacklist	False	mhermans

```
In [9]: len(set(df['author'])) # the number of author
```

```
Out[9]: 85838
```

2.Process the title(word vectorize)

```
In [10]: import nltk
nltk.download('punkt')
nltk.download('stopwords')
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
Out[10]: True
```

2.1build the corpus

```
In [0]: from nltk.stem.snowball import SnowballStemmer
stemmer = SnowballStemmer("english")
```

```
In [0]: # To get the stems of words in a sentence.
def tokenize_and_stem(text):
    # first tokenize by sentence, then by word to ensure that punctuation
    # is caught as it's own token
    tokens = [word for sent in nltk.sent_tokenize(text) for word in
nltk.word_tokenize(sent)]
    filtered_tokens = []
    for token in tokens:
        if re.search('[a-zA-Z]', token):
            filtered_tokens.append(token)

    stems = [stemmer.stem(t) for t in filtered_tokens]
    return stems

# To get the words themselves in a sentence.
def tokenize_only(text):
    # first tokenize by sentence, then by word to ensure that punctuation
    # is caught as it's own token
    tokens = [word for sent in nltk.sent_tokenize(text) for word in
nltk.word_tokenize(sent)]
    filtered_tokens = []
    for token in tokens:
        if re.search('[a-zA-Z]', token):
            filtered_tokens.append(token)
    return filtered_tokens
```

```
In [0]: #lowercase
title = df.title.str.lower()
```

```
In [0]: # Get full stems and tokens to build vocabulary
def tokenized_stemmed(title):
    totalvocab_stemmed = []
    totalvocab_tokenized = []
    for i in title:
        allwords_stemmed = tokenize_and_stem(i)
        totalvocab_stemmed.extend(allwords_stemmed)

        allwords_tokenized = tokenize_only(i)
        totalvocab_tokenized.extend(allwords_tokenized)
    return totalvocab_stemmed, totalvocab_tokenized
```

```
In [0]: totalvocab_stemmed_, totalvocab_tokenized_ = tokenized_stemmed(title)
```

```
In [16]: print(len(totalvocab_stemmed_))
```

7194561

```
In [0]: # pickle.dump((totalvocab_stemmed_, totalvocab_tokenized_), open("drive/My Drive/Eluvio/stem_token.pkl", "wb" ))
totalvocab_stemmed_, totalvocab_tokenized_ = pickle.load(open("drive/My Drive/Eluvio/stem_token.pkl", "rb" ))
```

```
In [0]: # Rule out repetitions of stem-token pairs
# totalvocab = zip(totalvocab_stemmed_, totalvocab_tokenized_)
# totalvocab = list(set(totalvocab))
# totalvocab_stemmed, totalvocab_tokenized = zip(*totalvocab)

# pickle.dump((totalvocab_stemmed, totalvocab_tokenized), open("drive/My Drive/Eluvio/stem_token.pkl", "wb" ))

totalvocab_stemmed, totalvocab_tokenized = pickle.load(open("drive/My Drive/Eluvio/stem_token.pkl", "rb" ))
```

```
In [19]: print(len(totalvocab_stemmed))
```

115041

```
In [0]: #stem-token vocabulary
# vocab_frame = pd.DataFrame({'words': totalvocab_tokenized}, index = totalvocab_stemmed)

# pickle.dump(vocab_frame, open('drive/My Drive/Eluvio/vocab_frame.pkl', 'wb'))

vocab_frame = pickle.load(open('drive/My Drive/Eluvio/vocab_frame.pkl', 'rb'))
```

```
In [0]: # Build stopwords set. Combine two common set.
import sklearn.feature_extraction.text as text
stopwords = nltk.corpus.stopwords.words('english')
my_stop_words = text.ENGLISH_STOP_WORDS.union(stopwords)
```

2.2 Tf-idf to vectorize text

```
In [22]: # tf-idf vectorizer
from sklearn.feature_extraction.text import TfidfVectorizer

tfidf_vectorizer = TfidfVectorizer(min_df =10**-3 ,analyzer = 'word',
max_features=len(set(totalvocab_stemmed)), stop_words=my_stop_words,
tokenizer=tokenize_and_stem, ngram_range=(1,3))

tfidf_matrix = tfidf_vectorizer.fit_transform(title)

print(tfidf_matrix.shape)

/usr/local/lib/python3.6/dist-packages/sklearn/feature_extraction/text.py:300: UserWarning: Your stop_words may be inconsistent with your preprocessing. Tokenizing the stop words generated tokens ['d', 's', 'abov', 'afterward', 'alon', 'alreadi', 'alway', 'ani', 'anoth', 'anyon', 'anyth', 'anywher', 'becam', 'becaus', 'becom', 'befor', 'besid', 'cri', 'describ', 'doe', 'dure', 'els', 'elsewhere', 'empti', 'everi', 'everyon', 'everyth', 'everywher', 'fifti', 'forti', 'henc', 'hereaft', 'herebi', 'howev', 'hundr', 'inde', 'mani', 'meanwhil', 'moreov', 'n't', 'need', 'nobodi', 'noon', 'noth', 'nowher', 'onc', 'onli', 'otherwis', 'ourselv', 'perhap', 'pleas', 'sever', 'sha', 'sinc', 'sincer', 'sixti', 'someon', 'someth', 'sometim', 'somewher', 'themselv', 'thenc', 'thereaft', 'therebi', 'therefor', 'togeth', 'twelv', 'twenti', 'veri', 'whatev', 'whenc', 'whenev', 'wherea', 'whereaft', 'wherebi', 'wherev', 'whi', 'wo', 'yourself'] not in stop_words.
  'stop_words.' % sorted(inconsistent))

(509236, 1814)
```

```
In [0]: # pickle.dump(tfidf_matrix, open("drive/My Drive/Eluvio/tfidf_matrix.pkl", "wb" ))
# pickle.dump(tfidf_vectorizer, open( "drive/My Drive/Eluvio/tfidf_vectorizer.pkl", "wb" ))

tfidf_matrix = pickle.load(open("drive/My Drive/Eluvio/tfidf_matrix.pkl", "rb" ))
# tfidf_vectorizer = pickle.load(open("drive/My Drive/Eluvio/tfidf_vectorizer.pkl", "rb" ))
```

```
In [25]: tfidf_matrix
```

```
Out[25]: <509236x1814 sparse matrix of type '<class 'numpy.float64'>'
with 3565328 stored elements in Compressed Sparse Row format>
```

Model

```
In [0]: thre = np.quantile(df['up_votes'], 0.8)
y = [1 if i > thre else 0 for i in df['up_votes']]
y = np.array(y)
X_train, X_test, y_train, y_test = train_test_split(tfidf_matrix, y
, test_size = 0.2, shuffle = True, random_state = 42)
```

MultinomialNB

```
In [27]: clf = MultinomialNB()
clf.fit(X_train, y_train)
```

```
Out[27]: MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)
```

```
In [28]: y_predict = clf.predict(X_test)
clf.score(X_test, y_test)
```

```
Out[28]: 0.8050624459979577
```

```
In [29]: print(classification_report(y_test, y_predict))
```

	precision	recall	f1-score	support
0	0.81	1.00	0.89	81988
1	0.56	0.00	0.00	19860
accuracy			0.81	101848
macro avg	0.68	0.50	0.45	101848
weighted avg	0.76	0.81	0.72	101848

LogisticRegression

```
In [0]: LR = LogisticRegression(C=1.0, penalty='l1', tol=0.01)
```

```
In [31]: LR.fit(X_train, y_train)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs'
  in 0.22. Specify a solver to silence this warning.
  FutureWarning)
```

```
Out[31]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=100,
                             multi_class='warn', n_jobs=None, penalty='l1',
                             random_state=None, solver='warn', tol=0.01, verbose=0,
                             warm_start=False)
```

```
In [32]: y_predict = LR.predict(X_test)
LR.score(X_test, y_test)
```

```
Out[32]: 0.8061719425025529
```

```
In [33]: print(classification_report(y_test, y_predict))
```

	precision	recall	f1-score	support
0	0.81	0.99	0.89	81988
1	0.55	0.04	0.07	19860
accuracy			0.81	101848
macro avg	0.68	0.51	0.48	101848
weighted avg	0.76	0.81	0.73	101848

GBDT


```
In [34]: gbdtd = GradientBoostingClassifier()
         gbdtd.fit(X_train, y_train)
```

```
Out[34]: GradientBoostingClassifier(criterion='friedman_mse', init=None,
                                   learning_rate=0.1, loss='deviance', max
                                   _depth=3,
                                   max_features=None, max_leaf_nodes=None,
                                   min_impurity_decrease=0.0, min_impurity
                                   _split=None,
                                   min_samples_leaf=1, min_samples_split=2
                                   ,
                                   min_weight_fraction_leaf=0.0, n_estimat
                                   ors=100,
                                   n_iter_no_change=None, presort='auto',
                                   random_state=None, subsample=1.0, tol=0
                                   .0001,
                                   validation_fraction=0.1, verbose=0,
                                   warm_start=False)
```

```
In [35]: y_predict = gbdtd.predict(X_test)
         gbdtd.score(X_test, y_test)
```

```
Out[35]: 0.8054257324640641
```

```
In [36]: print(classification_report(y_test, y_predict))
```

	precision	recall	f1-score	support
0	0.81	1.00	0.89	81988
1	0.77	0.00	0.01	19860
accuracy			0.81	101848
macro avg	0.79	0.50	0.45	101848
weighted avg	0.80	0.81	0.72	101848

Random Forest

```
In [37]: rfc = RandomForestClassifier(n_jobs = -1, max_features = 'sqrt', n_
estimators = 10, oob_score = True)
rfc.fit(X_train, y_train)
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:
460: UserWarning: Some inputs do not have OOB scores. This probabl
y means too few trees were used to compute any reliable oob estima
tes.
```

```
warn("Some inputs do not have OOB scores. ")
/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:
465: RuntimeWarning: invalid value encountered in true_divide
predictions[k].sum(axis=1)[: , np.newaxis])
```

```
Out[37]: RandomForestClassifier(bootstrap=True, class_weight=None, criterio
n='gini',
                                max_depth=None, max_features='sqrt', max_le
af_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_spl
it=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=
10, n_jobs=-1,
                                oob_score=True, random_state=None, verbose=
0,
                                warm_start=False)
```

```
In [38]: y_predict = rfc.predict(X_test)
rfc.score(X_test, y_test)
```

```
Out[38]: 0.7927107061503417
```

```
In [39]: print(classification_report(y_test, y_predict))
```

	precision	recall	f1-score	support
0	0.81	0.97	0.88	81988
1	0.30	0.05	0.08	19860
accuracy			0.79	101848
macro avg	0.56	0.51	0.48	101848
weighted avg	0.71	0.79	0.73	101848

XGB

```
In [0]: import xgboost as xgb
from xgboost.sklearn import XGBClassifier
```

```
In [0]: xgb = XGBClassifier(
        learning_rate =0.1,
        n_estimators=1000,
        max_depth=5,
        min_child_weight=1,
        gamma=0,
        subsample=0.8,
        colsample_bytree=0.8,
        objective= 'binary:logistic',
        nthread=4,
        scale_pos_weight=1,
        seed=27)
```

```
In [42]: xgb.fit(X_train, y_train)
```

```
Out[42]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=
1,
                colsample_bynode=1, colsample_bytree=0.8, gamma=0,
                learning_rate=0.1, max_delta_step=0, max_depth=5,
                min_child_weight=1, missing=None, n_estimators=1000,
n_jobs=1,
                nthread=4, objective='binary:logistic', random_state
=0,
                reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=
27,
                silent=None, subsample=0.8, verbosity=1)
```

```
In [0]: y_predict = xgb.predict(X_test)
```

```
In [44]: xgb.score(X_test, y_test)
```

```
Out[44]: 0.8062406723745189
```

```
In [45]: print(classification_report(y_test, y_predict))
```

	precision	recall	f1-score	support
0	0.81	0.99	0.89	81988
1	0.54	0.04	0.08	19860
accuracy			0.81	101848
macro avg	0.68	0.52	0.48	101848
weighted avg	0.76	0.81	0.73	101848

```
In [0]:
```